

AI readiness of the federal statistical system: host-level findings

AI readiness of the federal statistical system: host-level findings

Version. Snapshot cycle scan_2026-10-06_composite_c · version 2026-10-06_composite_c · released 2026-10-10; the commit this build was read from is recorded in data/index.json beside it. This document is a VIEW of data published beside it: the matrices as JSON and CSV, the per-check source appendix, and every Result quoted below with its value, its state and the artifact that generated it. The site index links all of them. **Licence.** The report and the data it is a view of are CC-BY-4.0 (LICENSE-DATA); the code that produced them is MIT (LICENSE). Documents under corpus/ are third-party works, retained as evidence under their own terms: neither licence in this repository applies to them. **Standing.** No later judgement of this cycle's evidence is on the event log: the snapshot is the current judgement of record.

Draft. 16 recognized bodies and three reference hosts, one cycle for the host checks and a second for twelve product checks, one client identity. Written from cc_tasks/2026-09-09_report_draft.md, revised to the cycle of 2026-09-10 under cc_tasks/2026-09-11_l0_report_cycle4_revision.md, re-snapshotted under cc_tasks/2026-09-19_resnapshot_rj4.md, again under cc_tasks/2026-10-06_absence_verdicts_rules.md, and onto a recollection under cc_tasks/2026-10-06_absence_verdicts_recollection_v2.md, and re-judged under cc_tasks/2026-10-07_seed_known_locations_and_split_discoverability.md. The snapshot is a declared composite, scan_2026-10-06_composite_c, of two cycles. The host checks, and every product check not listed next, are cycle scan_2026-09-10, the first over the complete frame, as re-judged in scan_2026-09-10_rj5. Twelve product checks were measured again on 2026-10-06, as cycle scan_2026-10-06_recollect, over the places each body declares for its API, its terms, its changelog and its inventories, and are shown as that evidence was re-judged on 2026-10-10 in scan_2026-10-06_recollect_rj1: A1, A2, A3, A9, B1, B3, B4, D1, D3, D4, F4 and G4. The re-judgement corrected three rules that had passed what they should not: a data catalogue served at an API's base is no longer read as an API description (A2), a licence name is no longer matched inside other words or inside a page's scripts (D1), and a web page is no longer taken for a downloadable file (A3). Every rule makes an absence a fail only over a complete search, and says error, naming what was not searched, everywhere else. Every number below is a registered Result quoted by name; nothing is typed into the prose.

What was measured, and over what

The frame is the recognized agencies and units of the federal statistical system: 16 bodies. They are taken from the Interagency Council on Statistical Policy's own charter, which enumerates its membership under an explicit flag legend rather than in prose. Recognition is a legal status conferred under CIPSEA 2018 (44 U.S.C. 3561, 3562), so the roster is a parse of an authoritative list and not a judgement about who counts. The charter and the Council's public About page disagree about 1 body, the Social Security Administration's research office. The charter marks it recognized and the About page does not, and the charter was named the authority before either document was read.

Three further hosts appear in the measurement and in none of its rates: data.gov, NIST and GSA. These are reference hosts. A federal catalog and two data portals are the nearest available control for what a machine-facing government site looks like when someone has tried. They are judged on the host-level checks only and enter no agency denominator, because placing a catalog beside a statistical agency above that level compares two different kinds of thing.

Counting surfaces rather than bodies, the target list carries 72 declared surfaces across 22 netlocs. Netlocs exceed bodies because each reference host declares a machine entry point on its own hostname.

Every surface is declared, by the roster or by the operator, and no rule selects one. Three selection rules were built for this frame and all three failed; the record of why is in the design decisions.

The client, and what it did

One identified client, one request per second per host, `robots.txt` obeyed, no forms, no logins, no query-string fuzzing. The user agent names the project and links to its source. The scanner has never retried a host under another identity, because a refusal is itself a measurement and a disguise would destroy it.

The first cycle issued 2684 requests across 35 netlocs, which is more netlocs than the target list names. The excess is not an accident and is discussed under what the matrix cannot see. It produced 2718 observations. The recollection issued 2074 requests under the same client and produced 6512 observations. The snapshot holds 1009 findings, and every finding re-derives byte for byte from its stored observations. The requests and the observations are each measured cycle's and carry its name; the findings are each judgement's, and the composite selects them by check without judging anything again.

The five checks

Each asks something a machine would need before it could use a body's statistics without a person in the loop.

A4 asks whether the host serves a `robots.txt` and whether it permits an identified, compliant client to read the page in question. A site with no such file has not refused; it has said nothing, and a client has to guess.

A5 asks whether the host offers a way to discover what it publishes, through a sitemap it declares itself or one of the newer machine-directed files. Discovery is the difference between a catalogue and a maze.

A10 asks whether a deep link behaves. A page that answers a nonsense URL with a cheerful page and a success status tells a machine that everything exists, which is worse than an honest absence.

A11-declared asks whether the host declares a machine layer anywhere in its own markup: an API, a bulk endpoint, a data catalogue.

A12 asks whether what the host declares and what it enforces agree, comparing the permission stated in `robots.txt` against the answer a compliant client actually receives on the same path. It is a candidate check, adopted by nobody, and it enters no fraction here.

The matrix

One row per body, one column per check, one cell per verdict. `pass` and `fail` are measurements. `error` means this scanner could not observe the surface, which is a fact about the scanner and its reception, never a fact about the product. The last column counts how many of the probes issued against the body's own page the host answered with a refusal status.

The four surface-judged checks are read from each body's home page, and the coherence check from its well-known set. Which surface a cell was measured on is recorded on every row of the machine-readable copy beside this one. It matters: a `robots.txt` that permits the front door can disallow a particular product, so a flagship page and a home page need not receive the same answer. This cycle found that disagreement on 12 cells across 6 bodies, and a single combined cell would have had to pick one without saying which.

Agency	A4	A5	A10	A11- declared	A12	Refused of probed
BEA	pass	fail	pass	pass	pass	1 of 129
BJS	pass	fail	pass	pass	pass	0 of 151
BLS	error	error	error	error	fail	34 of 38
BTS	error	error	error	error	fail	33 of 35
CENSUS	pass	pass	pass	pass	pass	6 of 166
DRSMSU	fail	fail	pass	fail	fail	0 of 267
EIA	pass	pass	pass	pass	pass	3 of 158
ERS	pass	fail	pass	pass	pass	1 of 112
NAHMSAPHIS	pass	fail	pass	pass	pass	1 of 202
NASS	pass	fail	pass	pass	pass	1 of 195
NCES	pass	fail	pass	pass	pass	1 of 81
NCHS	pass	fail	fail	fail	pass	1 of 37
NCSES	pass	pass	pass	pass	pass	0 of 97
ORES	error	error	error	error	fail	33 of 35
SAMHSACBHS	pass	fail	pass	pass	pass	2 of 927
SOI	pass	fail	pass	pass	pass	0 of 161
<i>this report</i>	-	-	-	-	-	no host yet; the instrument is turned on this report when it is published

Reading it by column

Serving robots.txt. 12 bodies of 13 observable ones serve a robots.txt that permits this client to read the page. The single failure serves no such file at all: not a refusal, an absence, and a machine meeting it has to assume rather than read. 3 bodies could not be observed on this check, and the same 3 recur in every column below.

Discovery. 3 of 13 offer a discoverable index of what they publish; 10 do not. Discovery is the weakest host-level result by a wide margin, and the consequence is direct: a machine that cannot enumerate what a site publishes cannot tell what it has missed. Every sitemap any host declared was followed, including ones declared on a neighbouring hostname, so none of these failures is an artefact of the scanner declining to look.

Deep links. 12 of 13 answer a deliberately invalid URL honestly. This is the best result on the page and it deserves less credit than it looks: it measures the absence of a specific pathology rather than the presence of a capability.

Declared machine layer. 11 of 13 declare somewhere in their own markup that a machine reader is expected. Declaring is not providing, and this check does not follow the declaration to see whether anything answers at the other end.

Declared against enforced. The candidate check finds 12 bodies coherent and 4 incoherent, of 16. None of the incoherent verdicts is the case the check is named for. That case is a declared layer that admits this client over an enforced layer that refuses it, and no body was observed in it. Three of the four are the Bureau of Labor Statistics, the Bureau of Transportation Statistics and the Social Security Administration’s research office. Their hosts refused the request for `/robots.txt` itself, so their declared layer could not be observed at all. The next section takes them up. The fourth, the Federal Reserve Board’s Microeconomic Surveys Unit, serves no `robots.txt` and so declares nothing for the enforced layer to agree or disagree with.

Uncertainty fields. This column has been withdrawn. **G1-D** asks whether a published figure carries, beside it and in machine-readable form, the fields that make it interpretable — a measure of uncertainty, a suppression flag, a reliability marker. The construct stands; the level was wrong. A body’s home page carries no estimate, so the check could not hold the property it measures there: across every body, every reference host and every cycle it returned 126 fail and 28 error on host-level surfaces and **not one** pass, while the same rule returned 104 pass on product surfaces in the same runs. A check that never passes on one kind of surface and passes 104 times on another is measuring which surface it was pointed at, not the publisher. It is withdrawn from the host-level instrument effective cycle 5 (this `params_hash` forward) (DD-066); the cycles that measured it keep every observation and every verdict on the append-only log, unchanged and un-re-judged, because it was not mis-scored — it was asked of the wrong surface. Where the property does live is a data product, and that is the January G1 pilot’s instrument, not this one’s.

The reference hosts

Agency	A4	A5	A10	A11-declared	A12	Refused of probed
GSA	pass	fail	pass	pass	pass	0 of 11
NIST	pass	fail	pass	pass	pass	0 of 11
data.gov	pass	fail	pass	pass	pass	0 of 12

All 3 reference hosts serve `robots.txt`, declare a machine layer, answer deep links honestly and are coherent between declaration and enforcement. All three fail discovery. They appear here for contrast and in no rate above.

Three bodies will not serve a compliant machine

3 bodies of 16 decline to answer a client that identifies itself, says where to complain about it, asks for no more than one page per second, and reads `robots.txt` first. They are the Bureau of Labor Statistics, the Bureau of Transportation Statistics and the Social Security Administration’s research office. Each returns a refusal status on effectively every request, `/robots.txt` included, so what these hosts declare to a machine has never been read. Their rows above read error throughout, which is the correct reading: this instrument did not find those sites wanting, it was not allowed to look.

The behaviour is neither a transient nor a sampling accident. It has now been recorded in 6 separate measurements, taken on four different days under two different user-agent strings. The count survives a change in how this project names errors: the earliest filed the refusal under a general client-error class, because the closed set of names had no member for a refusal until later. The number of bodies refusing has not moved, standing at 3 on the first look and the same on the most recent.

Two things follow, and only two. First, the coherence check fails on all three, and not because a declaration and a behaviour were seen to disagree: no declaration was seen. Every read of `/robots.txt`

on these hosts, in every cycle, was answered with HTTP 403, so the declared layer is not observable. RFC 9309 §2.3.1.3 calls a `robots.txt` answered with a status in the 400–499 range unavailable, and provides that then “the crawler MAY access any resources on the server”. The scanner was therefore permitted to go on, and the hosts refused the pages it asked for as they had refused the file. A client meeting them cannot learn the policy it is held to, because the one place the protocol puts that policy is refused as well. Second, no rate in this report describes them, and none can. They are in the frame, they are counted in the denominator of nothing, and the space they occupy in the matrix is the shape of what is not known.

What does not follow is any account of why. This scanner sees a status code. Bot management, a content delivery configuration, a deliberate policy and an unnoticed default all look identical from outside, and the only honest thing to report is the behaviour and its persistence. Distinguishing them needs either the operator’s own logs or a request made from a different vantage point, and both are named as future research rather than guessed at here.

In the previous cycle a fourth body could not be observed either, for an unrelated reason: its host timed out or closed the connection on most probes. It answered this time, and the three refusing bodies are now the whole of the unobserved column — which is why the error counts in the matrix are lower here than a reader of the last cycle would expect, and why they still are not zero.

What the products offer a machine

Everything above measures a host. This section measures products, and for the first time it covers the whole frame. Every one of the 16 bodies now has a flagship product declared for it — 16 agencies across 23 declared surfaces — and 0 bodies carry a host row and nothing else.

7 of those flagships were declared for this cycle, for the bodies that had none when the previous one ran, and each was verified robots-first before it entered the frame. 3 of them were answered with a refusal by the same hosts that refuse this client at every other door. They enter marked as refused rather than dropped: restricting the instrument to the agencies that permit it would make the frame a function of who answers. The rest came live in this cycle and are measured below.

Agency	Surface	A1	A2	A3	A6	A8	A9	B3	D1	D4	F4	B1	B2	B4	B5	D2	D3	G4
BEA	bea-	error	error	error	fail	fail	error	pass	error	error	error	error	fail	error	fail	fail	error	error
	flagship-1-interactive-data																	
BEA	bea-	error	error	error	fail	fail	error	pass	error	error	error	error	fail	error	fail	fail	error	error
	flagship-2-news-releases																	
BJS	bjs-	error	fail	error	fail	fail	fail	fail	error	fail	fail	fail	fail	fail	fail	fail	fail	fail
	flagship-1-data-by-topic																	
BJS	bjs-	error	fail	error	fail	fail	fail	fail	error	fail	fail	fail	fail	fail	fail	fail	fail	fail
	flagship-2-death-in-custody-reporting-act																	
flagship-1	www.dhs.gov/cpi/	error	error	error	error	error	error	error	error	error	error	error	error	error	error	error	error	error
flagship-1	www.dhs.gov/topics/national-transportation-statistics	error	error	error	error	error	error	error	error	error	error	error	error	error	error	error	error	error
CENSUS	census-	error	fail	error	fail	fail	pass	fail	error	error	error	error	fail	error	fail	fail	error	error
	flagship-1-surveys-programs																	
CENSUS	census-	error	fail	error	fail	fail	pass	fail	error	error	error	error	fail	error	fail	fail	error	error
	flagship-2-american-community-survey-acs																	
flagship-1	RSWSFederalassessmorgo/econres/scfindex.htm	error	error	pass	pass	pass	error	fail	error	fail	error	fail	fail	fail	fail	fail	fail	fail
EIA	eia-	error	error	error	error	error	error	error	error	fail	error	error	error	fail	error	fail	fail	fail
	flagship-1-open-data																	
ERS	ers-	pass	error	error	fail	fail	error	fail	error	error	error	error	fail	error	fail	fail	error	error
flagship-1	www.dhs.gov/flagship-1	error	error	fail	fail	error	fail	error	error	error	error	error	fail	error	fail	fail	error	error
flagship-1	www.dhs.gov/flagship-1	error	error	fail	fail	error	fail	error	error	error	error	error	fail	error	fail	fail	error	error
flagship-1	www.dhs.gov/flagship-1	error	error	fail	fail	error	fail	error	error	error	error	error	fail	error	fail	fail	error	error
flagship-1	www.dhs.gov/flagship-1	error	error	fail	fail	error	fail	error	error	error	error	error	fail	error	fail	fail	error	error
flagship-1	www.dhs.gov/flagship-1	error	error	fail	fail	error	fail	error	error	error	error	error	fail	error	fail	fail	error	error
flagship-1	www.dhs.gov/flagship-1	error	error	fail	fail	error	fail	error	error	error	error	error	fail	error	fail	fail	error	error
flagship-1	www.dhs.gov/flagship-1	error	error	fail	fail	error	fail	error	error	error	error	error	fail	error	fail	fail	error	error
flagship-1	www.dhs.gov/flagship-1	error	error	fail	fail	error	fail	error	error	error	error	error	fail	error	fail	fail	error	error
flagship-1	www.dhs.gov/flagship-1	error	error	fail	fail	error	fail	error	error	error	error	error	fail	error	fail	fail	error	error
flagship-1	www.dhs.gov/flagship-1	error	error	fail	fail	error	fail	error	error	error	error	error	fail	error	fail	fail	error	error
flagship-1	www.dhs.gov/flagship-1	error	error	fail	fail	error	fail	error	error	error	error	error	fail	error	fail	fail	error	error
flagship-1	www.dhs.gov/flagship-1	error	error	fail	fail	error	fail	error	error	error	error	error	fail	error	fail	fail	error	error
flagship-1	www.dhs.gov/flagship-1	error	error	fail	fail	error	fail	error	error	error	error	error	fail	error	fail	fail	error	error
flagship-1	www.dhs.gov/flagship-1	error	error	fail	fail	error	fail	error	error	error	error	error	fail	error	fail	fail	error	error
flagship-1	www.dhs.gov/flagship-1	error	error	fail	fail	error	fail	error	error	error	error	error	fail	error	fail	fail	error	error
flagship-1	www.dhs.gov/flagship-1	error	error	fail	fail	error	fail	error	error	error	error	error	fail	error	fail	fail	error	error
flagship-1	www.dhs.gov/flagship-1	error	error	fail	fail	error	fail	error	error	error	error	error	fail	error	fail	fail	error	error
flagship-1	www.dhs.gov/flagship-1	error	error	fail	fail	error	fail	error	error	error	error	error	fail	error	fail	fail	error	error
flagship-1	www.dhs.gov/flagship-1	error	error	fail	fail	error	fail	error	error	error	error	error	fail	error	fail	fail	error	error
flagship-1	www.dhs.gov/flagship-1	error	error	fail	fail	error	fail	error	error	error	error	error	fail	error	fail	fail	error	error
flagship-1	www.dhs.gov/flagship-1	error	error	fail	fail	error	fail	error	error	error	error	error	fail	error	fail	fail	error	error
flagship-1	www.dhs.gov/flagship-1	error	error	fail	fail	error	fail	error	error	error	error	error	fail	error	fail	fail	error	error
flagship-1	www.dhs.gov/flagship-1	error	error	fail	fail	error	fail	error	error	error	error	error	fail	error	fail	fail	error	error
flagship-1	www.dhs.gov/flagship-1	error	error	fail	fail	error	fail	error	error	error	error	error	fail	error	fail	fail	error	error
flagship-1	www.dhs.gov/flagship-1	error	error	fail	fail	error	fail	error	error	error	error	error	fail	error	fail	fail	error	error
flagship-1	www.dhs.gov/flagship-1	error	error	fail	fail	error	fail	error	error	error	error	error	fail	error	fail	fail	error	error
flagship-1	www.dhs.gov/flagship-1	error	error	fail	fail	error	fail	error	error	error	error	error	fail	error	fail	fail	error	error
flagship-1	www.dhs.gov/flagship-1	error	error	fail	fail	error	fail	error	error	error	error	error	fail	error	fail	fail	error	error
flagship-1	www.dhs.gov/flagship-1	error	error	fail	fail	error	fail	error	error	error	error	error	fail	error	fail	fail	error	error
flagship-1	www.dhs.gov/flagship-1	error	error	fail	fail	error	fail	error	error	error	error	error	fail	error	fail	fail	error	error
flagship-1	www.dhs.gov/flagship-1	error	error	fail	fail	error	fail	error	error	error	error	error	fail	error	fail	fail	error	error
flagship-1	www.dhs.gov/flagship-1	error	error	fail	fail	error	fail	error	error	error	error	error	fail	error	fail	fail	error	error
flagship-1	www.dhs.gov/flagship-1	error	error	fail	fail	error	fail	error	error	error	error	error	fail	error	fail	fail	error	error
flagship-1	www.dhs.gov/flagship-1	error	error	fail	fail	error	fail	error	error	error	error	error	fail	error	fail	fail	error	error
flagship-1	www.dhs.gov/flagship-1	error	error	fail	fail	error	fail	error	error	error	error	error	fail	error	fail	fail	error	error
flagship-1	www.dhs.gov/flagship-1	error	error	fail	fail	error	fail	error	error	error	error	error	fail	error	fail	fail	error	error
flagship-1	www.dhs.gov/flagship-1	error	error	fail	fail	error	fail	error	error	error	error	error	fail	error	fail	fail	error	error
flagship-1	www.dhs.gov/flagship-1	error	error	fail	fail	error	fail	error	error	error	error	error	fail	error	fail	fail	error	error
flagship-1	www.dhs.gov/flagship-1	error	error	fail	fail	error	fail	error	error	error	error	error	fail	error	fail	fail	error	error
flagship-1	www.dhs.gov/flagship-1	error	error	fail	fail	error	fail	error	error	error	error	error	fail	error	fail	fail	error	error
flagship-1	www.dhs.gov/flagship-1	error	error	fail	fail	error	fail	error	error	error	error	error	fail	error	fail	fail	error	error
flagship-1	www.dhs.gov/flagship-1	error	error	fail	fail	error	fail	error	error	error	error	error	fail	error	fail	fail	error	error
flagship-1	www.dhs.gov/flagship-1	error	error	fail	fail	error	fail	error	error	error	error	error	fail	error	fail	fail	error	error
flagship-1	www.dhs.gov/flagship-1	error	error	fail	fail	error	fail	error	error	error	error	error	fail	error	fail	fail	error	error
flagship-1	www.dhs.gov/flagship-1	error	error	fail	fail	error	fail	error	error	error	error	error	fail	error	fail	fail	error	error
flagship-1	www.dhs.gov/flagship-1	error	error	fail	fail	error	fail	error	error	error	error	error	fail	error	fail	fail	error	error
flagship-1	www.dhs.gov/flagship-1	error	error	fail	fail	error	fail	error	error	error	error	error	fail	error	fail	fail	error	error
flagship-1	www.dhs.gov/flagship-1	error	error	fail	fail	error	fail	error	error	error	error	error	fail	error	fail	fail	error	error
flagship-1	www.dhs.gov/flagship-1	error	error	fail	fail	error	fail	error	error	error	error	error	fail	error	fail	fail	error	error
flagship-1	www.dhs.gov/flagship-1	error	error	fail	fail	error	fail	error	error	error	error	error	fail	error	fail	fail	error	error
flagship-1	www.dhs.gov/flagship-1	error	error	fail	fail	error	fail	error	error	error	error	error	fail	error	fail	fail	error	error
flagship-1	www.dhs.gov/flagship-1	error	error	fail	fail	error	fail	error	error	error	error	error	fail	error	fail	fail	error	error
flagship-1	www.dhs.gov/flagship-1	error	error	fail	fail	error	fail	error	error	error	error	error	fail	error	fail	fail	error	error
flagship-1	www.dhs.gov/flagship-1	error	error	fail	fail	error	fail	error	error	error	error	error	fail	error	fail	fail	error	error
flagship-1	www.dhs.gov/flagship-1	error	error	fail	fail	error	fail	error	error	error	error	error	fail	error	fail	fail	error	error
flagship-1	www.dhs.gov/flagship-1	error	error	fail	fail	error	fail	error	error	error	error	error	fail	error	fail	fail	error	error
flagship-1	www.dhs.gov/flagship-1	error	error	fail	fail	error	fail	error	error	error	error	error	fail	error	fail	fail	error	error
flagship-1	www.dhs.gov/flagship-1	error	error	fail	fail	error	fail	error	error	error	error	error	fail	error	fail	fail	error	error
flagship-1	www.dhs.gov/flagship-1	error	error	fail	fail	error	fail	error	error	error	error	error	fail	error	fail	fail	error	error
flagship-1	www.dhs.gov/flagship-1	error	error	fail	fail	error	fail	error	error	error	error	error	fail	error	fail	fail	error	error
flagship-1	www.dhs.gov/flagship-1	error	error	fail	fail	error	fail	error	error	error	error	error	fail	error	fail	fail	error	error
flagship-1	www.dhs.gov/flagship-1	error	error	fail	fail	error	fail	error	error	error	error	error	fail	error	fail	fail	error	error
flagship-1	www.dhs.gov/flagship-1	error	error	fail	fail	error	fail	error	error	error	error	error	fail	error	fail	fail	error	error
flagship-1	www.dhs.gov/flagship-1	error	error	fail	fail	error	fail	error	error	error	error	error	fail	error	fail	fail	error	error
flagship-1	www.dhs.gov/flagship-1	error	error	fail	fail	error	fail	error	error	error	error	error	fail	error	fail	fail	error	error
flagship-1	www.dhs.gov/flagship-1	error	error	fail	fail	error	fail	error	error	error	error	error	fail	error	fail	fail	error	error
flagship-1	www.dhs.gov/flagship-1	error	error	fail	fail	error	fail	error	error	error	error	error	fail	error	fail	fail	error	error
flagship-1	www.dhs.gov/flagship-1	error	error	fail	fail	error	fail	error	error	error	error	error	fail	error	fail	fail	error	error
flagship-1	www.dhs.gov/flagship-1	error	error	fail	fail	error	fail	error	error	error	error	error	fail	error	fail	fail	error	error
flagship-1	www.dhs.gov/flagship-1	error	error	fail	fail	error	fail	error	error	error	error	error	fail	error	fail	fail	error	error
flagship-1	www.dhs.gov/flagship-1	error	error	fail	fail	error	fail	error	error	error	error	error	fail	error	fail	fail	error	error
flagship-1	www.dhs.gov/flagship-1	error	error	fail	fail	error	fail	error	error	error	error	error	fail	error	fail	fail	error	error
flagship-1	www.dhs.gov/flagship-1	error	error	fail	fail	error	fail	error	error	error	error	error	fail	error	fail	fail	error	error
flagship-1	www.dhs.gov/flagship-1	error	error	fail	fail	error	fail	error	error	error	error	error	fail	error	fail	fail	error	error
flagship-1	www.dhs.gov/flagship-1	error	error	fail	fail	error	fail	error	error	error	error	error	fail	error	fail	fail	error	error
flagship-1	www.dhs.gov/flagship-1	error	error	fail	fail	error	fail	error	error	error	error	error	fail	error	fail	fail	error	error
flagship-1	www.dhs.gov/flagship-1	error	error	fail	fail	error	fail	error	error	error	error	error	fail	error	fail	fail	error	error
flagship-1	www.dhs.gov/flagship-1	error	error	fail	fail	error	fail	error	error	error	error	error	fail	error	fail	fail	error	error
flagship-1	www.dhs.gov/flagship-1	error	error	fail	fail	error	fail	error	error	error	error	error	fail	error	fail	fail	error	error
flagship-1	www.dhs.gov/flagship-1	error	error	fail	fail	error	fail	error	error	error	error	error	fail	error				

Across the 23 declared surfaces, one product check still has no verdict at all, and that is the first thing to read off the matrix: a machine-readable licence (D1), error on every one of them. The terms pages the bodies declare were read, and none states a recognised licence in the text a reader sees; each was kept only in part, so the search for a licence on it is not complete and the check says error, naming how much was read. A documented API with its auth model and its rate limits (A2) now has a verdict where a body declares its API base, and the verdict is fail (the matrix's A2 column): an API is served at the declared base, and what it serves there is a catalogue of its datasets, not a description of the API. Everywhere else the body declares no API base and the check says error.

The other three checks that had no verdict now have some, because the search now reaches the places the indicator names. Structured data rather than a document is served on 4 of the 4 surfaces on which the question could be answered; on 19 more the product page links more of its own site than the scanner's request bound probed, and the links it did not reach may hold the data file. A machine-first entry point, a frontier check that enters no score, answers on 2 surfaces and is absent on 2 whose body declares its API base. A changelog is published on 8 surfaces whose body declares where its changes are recorded, and on every one of them it is a page for a person rather than a machine-readable feed; on 15 more no such place is declared, or the host refused to answer. Each error names what was not searched.

10 of the 17 product checks were measured and return not a single pass: variable-level metadata, concept definitions linked from the variables, data-quality attributes published as metadata rather than prose, terms that address model training and retrieval, published source lineage, an entry in a public inventory, the issuing authority recorded in that entry, and a machine-readable changelog. The last check at zero is asked of the body rather than of a surface: whether the same concept carries the same identifier across a body's products, and no body answers it yes. Four of these read the product's own record in a public inventory, and that record was found for no declared product. On 7 surfaces every inventory the body declares was read, its own host's, its department's and its data.gov organization's, and none names the product. On 16 more an inventory could not be read, because the department's or the body's own host refused this client or failed to answer, and the error names which. So the issuing authority and its programme, recorded as structured metadata in the product's catalogue record, are present on 0 of the 7 surfaces whose search was complete. The one product the earlier judgement found in an inventory, the Survey of Consumer Finances, is still in the Board's catalogue; its record now carries its landing page as an object instead of a URL, so no field the check reads names the product.

The other checks are not at zero. A bulk download of the whole product is linked from 6 of 6 surfaces on which the question could be answered. On 17 more it could not, because the links the scanner did not probe may hold the download. With structured data above, these are the pass rates on this page a reader would notice, and in both the denominator is why they should not be noticed for long. 3 of 19 serve their substantive content without requiring a browser to execute code. Structured markup describing the data and a resolvable pointer to the current vintage are present on 1 of 19 and 1 of 19 surfaces.

A low count at this denominator is not proof of universal absence, and the report will not let it be read as one. Structured markup describing the data is present on 1 of 19 surfaces. The upper bound of the ninety-five percent score interval on that rate is 0.246387, and that bound is the number a reader should carry rather than the count. Every check at zero has its own registered bound beside it, for the same reason. The bound quoted here is the score interval, computed once and registered rather than worked out in a sentence. The correct statement is that the true rate is unlikely to exceed that bound, not that it is what the count shows.

The interval is still wide, and the declaration list is no longer what makes it so. It is wide because one product is one surface, and the frame holds 19 observable surfaces on this check. 4 more answered with something this scanner could not read. A denominator that small is what a report on a system this size gets, and saying so is cheaper than pretending the bound is tighter.

What the matrix cannot see

Pages as a browser renders them. Every check here reads what the server sends. A site that assembles its content in the reader's browser looks empty to this scanner, and to any client that does not run code. That is a real finding about machine access and a poor description of what a person sees.

Whether any machine actually uses the data. Nothing here observes use. Access and uptake are different questions, and only the first is measured.

Anything about the three bodies that refused. Their cells are error and will stay that way under this method. No rate in this report includes them, and no reader should infer that their absence from a numerator means anything about their products.

Most of what "AI-ready" is taken to mean. The product section now covers every body in the frame, so the limit is no longer who is measured but what a check on a public surface can reach. Two published definitions say how much. NOAA's own order defines AI-ready data in 5 components, and this instrument measures 3 of them exactly — discoverable, machine-readable and machine-understandable, and access methods. Documentation it measures in part and quality not at all: the definition's word is *sufficient*, a judgement about the data against a use, and no check here opens the data. ESIP's AI-ready checklist has 58 assessable items and 9 of them are measured exactly here, most of those under data access. Neither gap is a hole in the framework, which names indicators for both; it is the distance between reading a surface and reading a dataset.

The scanner's own reach. The cycle contacted more netlocs than the target list names, which bounds what "the frame" means. Two bodies declare their sitemap on a neighbouring hostname and the scanner followed the declaration, which is what the discovery check is for. A closed contact list and a discovery check that reads what the host actually says cannot both hold; this cycle took the second and is reporting the choice rather than having made it in advance. One consequence is on the log: those requests went out without first reading the neighbouring host's own robots.txt, which this scanner's manners require. That is a defect in the collector rather than a finding about anyone, it is still open in this cycle, and it is queued.

The same declarations could have been declined instead, so the count of discovery failures attributable to a declaration the scanner refused to follow is registered at 0 for this cycle. That is zero, and it is registered rather than omitted: a caveat that is absent and a caveat measured to be empty read alike in prose and are not the same claim.

One failure this instrument cannot yet name. Exactly 1 observation in the cycle carries an error class of unknown: the map of failures does not have a name for what happened. It is counted and reported rather than folded into a neighbouring class, because an aggregate is only worth having if nothing was quietly swept into it.

What moved since the previous cycle, and what moved it

This time the evidence did not move, and three rules did. The figure below draws the snapshot beside the judgement it replaces, scan_2026-10-06_composite_b, the composite published on 2026-10-07. Both rest on the same stored observations. Three checks are marked as a rule change (A2, A3 and D1): each rule now reads the object as what it is rather than as bytes, so each can only have lost passes. A2 and D1 also change where their findings come from: the earlier composite withheld the recollection's verdicts on them and showed the 2026-09-10 cycle's, and this one shows the recollection's, re-judged. On every other check both points are the same findings. A check with no applicable denominator on a side has no point there, because every product surface is error on it.

the same evidence on both sides; generation 15 narrows what a pass may claim on A2, A3 and D1, and A2 and D1 now come from the 2026-10-06 recollection where composite B with




Figure 1: Pass rate per check, previous judgement beside this one

The figure draws no line and no arrow between the two points. A difference between two measurements is not a direction of travel, and two collections four weeks apart on federal publication schedules are not a rate of change. A check with no applicable denominator in a cycle is left blank rather than plotted at zero, because not measured is a reason and not a zero.

The judgements before this one moved the rule that reads the evidence, and each time in the same direction. A verdict about a product used to be allowed to rest on a probe the collector never made: a connection closed mid-request, a page a robots.txt forbade. Three successive corrections turned each of those into error. The fourth went further. An absence is a fail only over a search that reached everything the check names: every link on the product page, the agency's documented API, its department's inventory and data.gov's. Wherever the search stopped short, at the scanner's own request bound or at a location nobody declared, the verdict is error, and the error says which places were not searched. The recollection is the other half of that correction. The places were declared, each read off a page the body or data.gov publishes, and searched; a check now says fail where the whole declared search found nothing, and still says error where a host refused to be searched.

The bulk-download check shows both halves. Under the first re-judgement of the first cycle it was answered over 19 declared flagship surfaces. The fourth correction cut that sharply: on some surfaces the scanner was forbidden to look, and on the others the page's links ran past the request bound before the download could be reached, and both kinds had been counted as products that offer none. The recollection ranks a page's links by how much each looks like data before the bound applies, and the question is now answered over 6. The upper bound of the ninety-five percent interval on the rate went

from 0.433343 under the first re-judgement to 1 now. A larger denominator is a stronger claim only because the search behind it reached further, and an instrument that stops scoring what it did not look at has to earn every surface it adds back.

What this cannot answer, and what would

Why three bodies refuse an identified client. A status code cannot distinguish a policy from a default. Asking the operators, or reading their own access logs, is the only method that separates them, and it is the highest-value open question in this report.

Whether a third party sees what we see. Every observation here comes from one vantage point under one identity. A request from elsewhere, or under a common crawler's name, would establish whether the refusals are about this client or about machines in general.

What the products offer beyond one page each. The declaration list is complete and this cycle ran under it, so the open question is no longer coverage but depth: a flagship landing page is one surface, and a body's data lives behind it. Whether these checks answer the same way on a product's download, its API and its documentation pages is not known from a landing page, and asking it needs a declared surface per product rather than per body.

Whether the contact bound should be closed or open. Following a sitemap declared on a neighbouring host is either correct discovery or scope creep, and this cycle did it without having decided. One resolution bounds such following to the frame and records an off-frame declaration as an observation in its own right. The other states the bound as the frame plus whatever the frame's own hosts declare.

This report, measured by its own checks

The last row of the matrix is this report's own host, and it carries no verdicts. The report has no host: it exists as a file in a repository, so it serves no `robots.txt`, declares no machine layer and answers no deep links, and the intended home is GitHub Pages under this repository. Scoring it now would mean inventing answers for a site that does not exist, so the row stays empty until the instrument can be turned on this report's own output.

What the design calls for once it is published is on the record and is not vague. It wants a machine-readable copy of the matrix beside the readable one, and both already exist as the comma-separated and structured files accompanying this draft. It also wants a declared licence, a catalogue entry, markup identifying the matrix as a dataset, and a `robots.txt` permitting an identified client to read all of it. The instrument is meant to be turned on its own output. Until it can be, this section says so rather than reporting a row nobody measured.

Method appendix

Where the numbers live. Every value in this report is a registered Result in the project's artifact graph, quoted by name and resolved at build time. Nothing was typed. The build refuses to write this file if a single reference fails to resolve, and a lint over the built text refuses any numeral in prose that did not come through a reference. The matrices are rendered from the cycle's own data rather than transcribed. Each cell in the machine-readable copies names the Finding identity it came from, so any cell can be walked back to its observations, and those to the retained response bodies.

Client identity. `ai-readiness-kg-scanner/0.2` (`+https://github.com/brockwebb/ai-readiness-kg`). One identity, declared in parameters and read from there by everything that reports it. No host was ever retried under another.

Manners. RFC 9309 obeyed: `robots.txt` read before fetching, disallowed paths not fetched and recorded as a decision rather than as a failure. One request per second per host, one worker, no forms, no logins, no query-string fuzzing. Retries only on 429 and 503.

Verdicts. `pass` and `fail` are measurements. `error` means the collector could not observe and is excluded from every denominator; it is never a product failure. `n.a.` means there was nothing of that kind to check. `not declared` means no product has been named.

Cells read from a parent organization's host. Some statistical units publish on a host whose root belongs to an organization above them, so that host's `robots.txt`, `/data.json` and well-known files answer for that organization, and a verdict on them is not a finding about the unit. The machine-readable matrices keep those cells, with their verdicts and Findings, and mark them `parent_host` in each row's marks. The per-check rates in this report still count them, because those rates describe hosts and surfaces. The project's scores and ranks, which this report does not print, leave them out. The bodies, the cells and the roster's reason for each, generated from the roster:

Body	Host	Whose host	Cells marked on parent_host on this cycle	Why the roster says so
DRSMSU	www.federalreserve.gov	Board of Governors of the Federal Reserve System	A4, A5, A11-declared, A12, D4, B1, B4, D2, D3, G4	the frame's home surface for the body is https://www.federalreserve.gov/econres/rsmecs-staff.htm, a section of www.federalreserve.gov rather than its root
NAHMSAPHIS	www.aphis.usda.gov	Department of Agriculture	A4, A5, A11-declared, A12, D4, B1, B4, D2, D3, G4	the frame's home surface for the body is https://www.aphis.usda.gov/aphis/ourfocus/animalhealth/monitoring-and-surveillance/nahms, a section of www.aphis.usda.gov rather than its root
NCHS	www.cdc.gov	Centers for Disease Control and Prevention	A4, A5, A11-declared, A12, D4 (2), B1 (2), B4 (2), D2 (2), D3 (2), G4 (2)	the cycle-1 roster carries host_shared_with: Centers for Disease Control and Prevention (assessment/harness/scan/targets.yaml); the frame's home surface for the body is https://www.cdc.gov/nchs/index.htm, a section of www.cdc.gov rather than its root
ORES	www.ssa.gov	Social Security Administration	A4, A5, A11-declared, A12, D4, B1, B4, D2, D3, G4	the .gov registry attributes ssa.gov to Social Security Administration,
SAMHSAICBHS	www.samhsa.gov	Substance Abuse and Mental Health Services Administration	A4, A5, A11-declared, A12, D4, B1 (2), B4 (2), D2 (2), D3 (2), G4 (2)	the roster's the parent's home surface for the department, body is https://www.samhsa.gov/suborganization/samhsa/ and www.ssa.gov is that domain

Frontier checks. The machine-first entry point check (A9) is a frontier mechanism, dated by its indicator. It is on the product matrix with its verdicts and is marked `frontier`; a frontier mechanism enters no score.

Intervals. Wilson score intervals, computed by the project's own rollup module and registered per check. Wilson (1927) for the interval. Brown, Cai and DasGupta (2001) and Newcombe (1998) for why a score interval rather than a normal approximation at proportions near zero or one. Hanley and Lippman-Hand (1983) for the rule of three at zero events.

Controls. Seven local fixture servers were scanned in this cycle alongside the real hosts. One passes every check and one fails every check. One refuses an identified client and one resets the connection. One answers normally but kills the invalid-route probe, and one answers every GET while resetting every HEAD. The seventh declares its sitemap on a second hostname of the same site, which is the case the discovery check meets on two real hosts; what it pins is an order, that the sibling's own `robots.txt` is read before anything else is asked of it. Their expected verdicts are derived from what each collector dispatches, not written by hand: a hand-written expectation was wrong once.

Rules that judged this cycle.

Check	Rule that judged this cycle
A1	RULE-A1-v5
A10	RULE-A10-v3
A11-declared	RULE-A11-declared-v2
A12	RULE-A12-v3
A2	RULE-A2-v5
A3	RULE-A3-v8
A4	RULE-A4-v1
A5	RULE-A5-v2
A6	RULE-A6-v2
A8	RULE-A8-v4
A9	RULE-A9-v2
B1	RULE-B1-v3
B2	RULE-B2-v1
B3	RULE-B3-v4
B4	RULE-B4-v2
B5	RULE-B5-v1
D1	RULE-D1-v5
D2	RULE-D2-v1
D3	RULE-D3-v2
D4	RULE-D4-v4
F4	RULE-F4-v4
G1-D	RULE-G1-D-v1
G4	RULE-G4-v2

Requests issued, per netloc. Summed over the two collections the snapshot rests on: the first cycle, which measured every check, and the recollection, which measured the twelve checks listed in the frame section.

Netloc	Requests
api.census.gov	9
api.eia.gov	7
api.ojp.gov	7
apps.bea.gov	34
bjs.ojp.gov	325
catalog.data.gov	60
data.bls.gov	2
data.census.gov	11
data.gov	2
data.nist.gov	13
datatools.samhsa.gov	3
gis.cdc.gov	4
hhs.gov	5
ir.eia.gov	4
nces.ed.gov	223
ncses.nsf.gov	330
open.gsa.gov	13
quickstats.nass.usda.gov	2
sa.www4.irs.gov	17
samhsa.gov	3
tools.bea.gov	10
tools.cdc.gov	5
wonder.cdc.gov	4
www.aphis.usda.gov	214
www.bea.gov	353
www.bls.gov	118
www.bts.gov	114
www.cdc.gov	369
www.census.gov	430
www.commerce.gov	9
www.data.gov	15
www.dol.gov	3
www.ed.gov	3
www.eia.gov	253
www.energy.gov	4
www.fda.gov	408
www.fishbase.org	408

Sources per check.

One row per check and admitted source: the check as the report names it, the framework indicator it measures, the source as the corpus manifest records it, its document id in the manifest, and the locator inside the source as the indicator's evidence cell states it. Generated from the graph; nothing in this table is typed.

Check	Indicator	Source	doc_id	Locator
A4	A4 Crawler/ agent access	· Anthropic. <i>Does Anthropic crawl data from the web, and how can site owners block the crawler? (Anthropic support)</i> . https://support.anthropic.com/en/articles/8896518-does-anthropic-crawl-data-from-the-web-and-how-can-site-owners-block-the-crawler .	anthropic-crawler-support-article	
A4	A4 Crawler/ agent access	· Cloudflare. <i>Cloudflare AI Crawl Control: Manage AI crawlers</i> . 2026. https://developers.cloudflare.com/ai-crawl-control/features/manage-ai-crawlers/ .	cloudflare-ai-crawl-control-manage-crawlers	
A4	A4 Crawler/ agent access	· Google Search Central. <i>Google Search Central: Introduction to robots.txt</i> . 2025. https://developers.google.com/search/docs/crawling-indexing/robots/intro .	google-robots-txt-intro	
A4	A4 Crawler/ agent access	· OpenAI. <i>OpenAI: Overview of OpenAI crawlers (developers.openai.com/api/docs/bots)</i> . https://platform.openai.com/docs/bots .	openai-crawlers-bots	
A4	A4 Crawler/ agent access	· Perplexity AI. <i>Perplexity crawlers (docs.perplexity.ai)</i> . https://docs.perplexity.ai/guides/bots .	perplexity-crawlers	
A4	A4 Crawler/ agent access	· M. Koster; G. Illyes; H. Zeller; L. Sassman; IETF. <i>RFC 9309: Robots Exclusion Protocol</i> . 2022. https://www.rfc-editor.org/rfc/rfc9309 .	rfc-9309-robots-exclusion-protocol	
AA5	· Discoverability surface	· Jeremy Howard / Answer.AI. <i>The /llms.txt file (llmstxt.org)</i> . 2026. https://llmstxt.org/ .	llmstxt-proposal	
AA5	· Discoverability surface	· sitemaps.org. <i>Sitemaps XML format (sitemaps.org protocol)</i> . 2022. https://www.sitemaps.org/protocol.html .	sitemaps-protocol	

census.gov/data/2020/decade/c

observed)

Files beside this report. `scan_matrix_tierA_2026-10-06_composite_c.csv` and `.json`, the host-level matrix; `scan_matrix_tierC_2026-10-06_composite_c.*`, the reference hosts; `scan_matrix_product_2026-10-06_composite_c.*`, the product matrix. Every row carries its Finding identities, and each JSON file names the two cycles under `composed_of`.

Provenance. A declared composite, `scan_2026-10-06_composite_c` (`state/scan_2026-10-06_composite_c.json`, `composed_of`). The host checks and every product check but twelve are cycle `scan_2026-09-10`, parameter hash `4e0a92ba19ab769bb98b3a4a0c68640f465a04eaaec4aa6f2f0f41dc75c0df`, judged as `scan_2026-09-10_rj5`: the same stored observations under the rules current on 2026-10-06. The twelve (A1, A2, A3, A9, B1, B3, B4, D1, D3, D4, F4, G4) are cycle `scan_2026-10-06_recollect`, parameter hash `2e56a8815bc22a2708f153978a8231ccfa1647027f85f27c5979beecfa1b50e9`, collected on 2026-10-06 over the API, terms, changelog and inventory locations each body declares (`targets.yaml` declared `locations`, every entry citing the page it was read from), and judged as `scan_2026-10-06_recollect_rj1` on 2026-10-10, parameter hash `bcea2d92c74e53f8d9e364c7eee0ab143675ddba6574553c1473197d18d926aa`, under generation 15 of the rules, which corrects three false passes on A2, D1 and A3. The parameters moved between the collection and the judgement only in blocks the twelve legs' rules do not read (the existence and discoverability blocks, and the discoverability candidate's entry under the link probe), so no other verdict moved. Every superseded judgement stays registered under its own name. The event log is the source of truth; the graph and the matrices are projections of it and are rebuilt by replay. Design decisions DD-059 (the frame and the tier separation), DD-060 (one client identity), DD-061 (control tables derived from collector dispatch) and DD-064 (forbidden to look is blindness, outside the product is scope) govern what this report may say.